

Project 6 - Ozzy Simpson

Overall Performance of Optimized Code:

Performance improvements by time in milliseconds

| Dimension | naive_rotate | my_rotate | Speedup | naive_smooth | my_smooth | Speedup |
|-----------|--------------|-----------|---------|--------------|-----------|---------|
| 256 | | | | 6799 | 1940 | 350.46% |
| 512 | 1020 | 753 | 135.46% | 30430 | 7879 | 386.22% |
| 1024 | 8149 | 3280 | 248.45% | 155057 | 34212 | 453.22% |
| 2048 | 54719 | 45690 | 119.76% | 834334 | 134984 | 618.10% |
| 4096 | 627518 | 158877 | 394.97% | | | |
| | | Average: | 224.66% | | Average: | 452.00% |

Performance improvements by time in milliseconds

| Dimension | naive_rotate | my_rotate | Speedup | naive_smooth | my_smooth | Speedup |
|-----------|--------------|-----------|---------|--------------|-----------|---------|
| 256 | | | | 13630951 | 3881243 | 351.20% |
| 512 | 2069471 | 1463788 | 141.38% | 60866341 | 15760623 | 386.19% |
| 1024 | 14532535 | 6999360 | 207.63% | 311164828 | 68428119 | 454.73% |
| 2048 | 120437954 | 66039451 | 182.37% | 1671769911 | 270061108 | 619.03% |
| 4096 | 860793384 | 306333113 | 281.00% | | | |
| | | Average: | 203.09% | | Average: | 452.79% |

my_rotate():

Function Inlining

While it is a macro (and not a function), we can inline the RIDX calculation. This will become helpful later on when we do code motion.

| | |
|---|--|
| <i>Before:</i> <code>dst[RIDX(dim-1-j, i, dim)] = src[RIDX(i, j, dim)];</code> | <i>After:</i> <code>dst[(dim-1-j) * dim + i] = src[i * dim + j];</code> |
|---|--|

This change resulted in a slight increase in performance because the calculations are performed inline and the macro doesn't have to be called; on average, it increased the performance by 3%:

| Dimension | naive_rotate | my_rotate | Speedup |
|-----------|--------------|-----------|---------|
| 512 | 1084 | 1029 | 105.34% |
| 1024 | 7239 | 6904 | 104.85% |
| 2048 | 51228 | 51614 | 99.25% |
| 4096 | 549993 | 534463 | 102.91% |
| | | Average: | 103.09% |

Dataflow Analysis and Optimizations

There were not many data flow optimizations that could be performed. The original rotate function does not include any dead code and there were not any opportunities to do constant folding or copy propagation. There are some common expressions that can be simplified, but this will be done in the code motion section below because we'll move common expressions outside of some loops.

Code Motion and Strength Reduction

We can notice that `RIDX()` is called with `dim-1-j` as the first argument every iteration of the inner loop, but `dim-1-j` only changes with each iteration of the outer loop. We can simplify this by moving this invariant computation outside of the inner loop, storing it in a temporary variable, and then using that as one of the arguments for the `dst` index calculation, like this:

| | |
|---|---|
| <i>Before:</i> <pre>for (j = 0; j < dim; j++) for (i = 0; i < dim; i++) dst[RIDX(dim-1-j, i, dim)] = src[RIDX(i, j, dim)];</pre> | <i>After:</i> <pre>int in; for (j = 0; j < dim; j++) { in = dim - 1 - j; for (i = 0; i < dim; i++) dst[RIDX(in, i, dim)] = src[RIDX(i, j, dim)]; }</pre> |
|---|---|

This change (alone) resulted in a more significant increase in performance, which is to be expected since we're calculating *in* fewer times (*dim* fewer times, to be exact). On average, this increased performance by 7%:

| Dimension | naive_rotate | my_rotate | Speedup |
|-----------|--------------|-----------|---------|
| 512 | 1011 | 961 | 105.20% |
| 1024 | 7940 | 6438 | 123.33% |
| 2048 | 62486 | 61511 | 101.59% |
| 4096 | 504535 | 507486 | 99.42% |
| | | Average: | 107.38% |

Furthermore, we can combine this change with the previous change and further eliminate common subexpressions to calculate part of the index completely out of the loops and as much as we can of the rest just in the outer loop, such that the for loops have been changed as follows:

| | |
|---|--|
| <i>Before:</i> <pre>for (j = 0; j < dim; j++) for (i = 0; i < dim; i++) dst[RIDX(dim-1-j, i, dim)] = src[RIDX(i, j, dim)];</pre> | <i>After:</i> <pre>int in; in = dim * dim - dim; for (j = 0; j < dim; j++) { for (i = 0; i < dim; i++) { dst[in + i] = src[i * dim + j]; } in -= dim; }</pre> |
|---|--|

This is possible because $(dim-1-j) * dim + i = dim * dim - dim - j * dim + i$. Anything not attached to i and j can be brought fully out of the loops, and any terms with just j can be put in the outer loop (but, can be simplified using strength reduction such that we subtract a dim at the end of each iteration of the outer loop rather than using the more costly multiplication).

This change gave the below performance results (average speedup from the original of 12%). This is because we calculate the more computationally expensive parts only once outside of the loops and just perform basic subtraction in the outer loop.

| Dimension | naive_rotate | my_rotate | Speedup |
|-----------|--------------|-----------|---------|
| 512 | 1068 | 1003 | 106.48% |
| 1024 | 7449 | 6711 | 111.00% |
| 2048 | 63086 | 50613 | 124.64% |
| 4096 | 436503 | 407986 | 106.99% |
| | | Average: | 112.28% |

Improving memory access/locality

There are four main ways in which we can improve memory access/locality: merging arrays, loop interchange, loop fusion, and blocking. Merging arrays, loop interchange, and loop fusion is not possible for this function (we can't have a single array, changing the nesting of loops doesn't change the fact that we need to use both row-major and column-major order, and there aren't loops to fuse together), but blocking can be used to improve performance since we cannot avoid accessing data in column-major order. Since the image dimensions are a multiple of 32, we can use blocks of 32 to improve the locality of our memory accesses. In each block, our accesses to `dst` are in order such that after we access `dst[0]` we access `dst[1]`; for our access to `src`, we're accessing `src[0]`, `src[32]`, `src[64]`, ... and then `src[1]`, `src[33]`, `src[65]`, ... instead of something like `src[0]`, `src[512]`, `src[1024]`, This means we have a higher cache hit rate because it's more likely the elements we're accessing are in the cache.

The modified code, taking advantage of (similar) improvements we performed in previous sections, now looks like:

| | |
|---|---|
| <i>Before:</i> <pre>for (j = 0; j < dim; j++) for (i = 0; i < dim; i++) dst[RIDX(dim-1-j, i, dim)] = src[RIDX(i, j, dim)];</pre> | <i>After:</i> <pre>int dim2, k; dim2 = dim * dim; dst += dim2 - dim; for (j = 0; j < dim; j += 32) { for (i = 0; i < dim; i++) { for (k = 0; k < 31; k++) { dst[k] = *src; src += dim; } dst[k] = *src; dst -= dim; src -= dim*31 - 1; } dst += 32 + dim2; src += dim * 31; }</pre> |
|---|---|

We can see that this change made a more significant change to the performance of the `my_rotate()` function, with an average improvement of 101%::

| Dimension | naive_rotate | my_rotate | Speedup |
|-------------|--------------|-----------|---------|
| 512 | 1016 | 691 | 147.03% |
| 1024 | 7827 | 4750 | 164.78% |
| 2048 | 63091 | 35682 | 176.81% |
| 4096 | 585942 | 185253 | 316.29% |
| | | Average: | 201.23% |

With a block size of 16, though, the performance is further improved (to a ~125% increase), which indicates that we are increasing memory locality and reducing cache misses over the block size of 32 (likely due to the size of the cache):

| Dimension | naive_rotate | my_rotate | Speedup |
|-------------|--------------|-----------|---------|
| 512 | 1020 | 753 | 135.46% |
| 1024 | 8149 | 3280 | 248.45% |
| 2048 | 54719 | 45690 | 119.76% |
| 4096 | 627518 | 158877 | 394.97% |
| | | Average: | 224.66% |

my_smooth():

Function Inlining

There are many opportunities to inline functions with the smoothing function. We can inline the RIDX() macro, the avg() function, and the functions that the avg() function calls (initialize_pixel_sum(), maximum(), minimum(), accumulate_sum(), and assign_sum_to_pixel()) which will improve performance by avoiding the overhead that comes from functions (activation records, return instructions, etc.). The code now looks like this:

Before:

```
for (j = 0; j < dim; j++)
    for (i = 0; i < dim; i++)
        dst[RIDX(i, j, dim)] = avg(dim, i, j, src);
```

After:

```
int ii, jj;
pixel_sum sum;
pixel current_pixel;
for (j = 0; j < dim; j++) {
    for (i = 0; i < dim; i++) {
        sum.red = sum.green = sum.blue = 0;
        sum.num = 0;
        for (ii = ((i - 1) > 0 ? i - 1 : 0); ii <= ((i + 1) < (dim - 1) ? (i + 1) : (dim - 1)); ii++)
            for (jj = ((j - 1) > 0 ? j - 1 : 0); jj <= ((j + 1) < (dim - 1) ? (j + 1) : (dim - 1)); jj++) {
                sum.red += (int)src[(ii * dim + jj)].red;
                sum.green += (int)src[(ii * dim + jj)].green;
                sum.blue += (int)src[(ii * dim + jj)].blue;
                sum.num++;
            }

        current_pixel.red = (unsigned short)(sum.red / sum.num);
        current_pixel.green = (unsigned short)(sum.green / sum.num);
        current_pixel.blue = (unsigned short)(sum.blue / sum.num);
        dst[(i * dim + j)] = current_pixel;
    }
}
```

The performance improvements from just inlining all of the functions is significant, with an average speedup of nearly 56%:

| Dimension | naive_smooth | my_smooth | Speedup |
|-----------|--------------|-----------|---------|
| 256 | 7320 | 4687 | 156.18% |
| 512 | 32177 | 20004 | 160.85% |
| 1024 | 155367 | 99789 | 155.70% |
| 2048 | 820113 | 548003 | 149.65% |
| | | Average: | 155.59% |

Dataflow Analysis and Optimizations, Code Motion

Using the code from the previous optimization as a starting point, we can perform some dataflow analysis and do constant folding and eliminate common subexpressions:

- We notice that the bound for the ii for loop is calculated each time the loop iterates; we can improve performance by calculating it once before the loop and storing the value in a temp variable so that it doesn't need to be calculated each time ii is incremented. We can do the same for the bound for the for loop that uses jj as an iterator.
- We also notice that $ii * dim + jj$ is calculated 3 times each time the jj loop iterates; we can calculate $ii * dim$ outside of the jj loop since it doesn't rely on jj and then just add jj to that value inside of the jj loop.
- We also see that $i + 1$ and $j + 1$ are calculated more than once per loop of their respective loops, so we store those calculated values in temp variables to avoid re-calculating them and improve performance. The same can be said for $i - 1$ and $j - 1$.
- $dim - 1$ is also calculated many times, so we'll take that out of all of the for loops entirely since it does not depend on any iterators and only calculate it once per function call.

After:

```

int ii, jj, iidim, iidimjj, iiEnd, jjEnd, dimMinusOne, iPlusOne, jPlusOne, iMinusOne, jMinusOne;
pixel_sum sum;
pixel current_pixel;
dimMinusOne = dim - 1;

for (j = 0; j < dim; j++) {
    jPlusOne = j + 1;
    jMinusOne = j - 1;
    for (i = 0; i < dim; i++) {
        iPlusOne = i + 1;
        iMinusOne = i - 1;
        sum.red = sum.green = sum.blue = 0;
        sum.num = 0;
        iiEnd = (iPlusOne < dimMinusOne ? iPlusOne : dimMinusOne);
        for (ii = (iMinusOne > 0 ? iMinusOne : 0); ii <= iiEnd; ii++) {
            iidim = ii * dim;
            jjEnd = (jPlusOne < dimMinusOne ? jPlusOne : dimMinusOne);
            for (jj = (jMinusOne > 0 ? jMinusOne : 0); jj <= jjEnd; jj++) {
                iidimjj = iidim + jj;
                sum.red += (int)src[iidimjj].red;
                sum.green += (int)src[iidimjj].green;
                sum.blue += (int)src[iidimjj].blue;
                sum.num++;
            }
        }

        current_pixel.red = (unsigned short)(sum.red / sum.num);
        current_pixel.green = (unsigned short)(sum.green / sum.num);
        current_pixel.blue = (unsigned short)(sum.blue / sum.num);
        dst[(i * dim + j)] = current_pixel;
    }
}

```

These changes net us an average speedup from the original function of ~74%:

| Dimension | naive_smooth | my_smooth | Speedup |
|-------------|--------------|-----------|---------|
| 256 | 7368 | 3758 | 196.06% |
| 512 | 31024 | 18107 | 171.34% |
| 1024 | 158816 | 88126 | 180.21% |
| 2048 | 824952 | 549482 | 150.13% |
| | | Average: | 174.44% |

Improving memory access/locality

It is not possible to merge arrays with this function since we don't have two arrays to merge. It is also not possible to benefit from loop fusion since there are no loops that have the same looping with overlapping variables.

The best we can do with memory blocking is using row-major access for the dst pixels. We notice that the smooth function iterates through every pixel and smooths each one. However, the order of nesting of the for loops results in the image being smoothed column-by-column, rather than row-by-row, which is unnecessary. Since C uses row-major order, this is computationally heavy and can result in cache misses and more memory accesses. We can use loop interchange to build on our last improvements and improve locality and reduce cache misses. This also gives us the best form of memory blocking, since we're accessing in row-major order. (We also can now calculate $i * \text{dim}$ outside of the j loop, making that calculation run significantly fewer times—from dim^2 times to just dim times.)

After:

```
int ii, jj, iDim, iidimjj, iiEnd, jjEnd, dimMinusOne, iPlusOne, jPlusOne, iMinusOne, jMinusOne;
pixel_sum sum;
pixel_current_pixel;
dimMinusOne = dim - 1;

for (i = 0; i < dim; i++) {
    iPlusOne = i + 1;
    iMinusOne = i - 1;
    iDim = i * dim;
    for (j = 0; j < dim; j++) {
        jPlusOne = j + 1;
        jMinusOne = j - 1;
        sum.red = sum.green = sum.blue = 0;
        sum.num = 0;
        iiEnd = (iPlusOne < dimMinusOne ? iPlusOne : dimMinusOne);
        for (ii = (iMinusOne > 0 ? iMinusOne : 0); ii <= iiEnd; ii++) {
            iidimjj = ii * dim;
            jjEnd = (jPlusOne < dimMinusOne ? jPlusOne : dimMinusOne);
            for (jj = (jMinusOne > 0 ? jMinusOne : 0); jj <= jjEnd; jj++) {
                iidimjj = iDim + jj;
                sum.red += (int)src[iidimjj].red;
                sum.green += (int)src[iidimjj].green;
                sum.blue += (int)src[iidimjj].blue;
                sum.num++;
            }
        }

        current_pixel.red = (unsigned short)(sum.red / sum.num);
        current_pixel.green = (unsigned short)(sum.green / sum.num);
        current_pixel.blue = (unsigned short)(sum.blue / sum.num);
        dst[(iDim + j)] = current_pixel;
    }
}
```

By swapping the order of the i and j loops, we now smooth row-by-row, which, as predicted, increases performance significantly (average of 135% from the original function):

| Dimension | naive_smooth | my_smooth | Speedup |
|-----------|--------------|-----------|---------|
| 256 | 7415 | 4014 | 184.73% |
| 512 | 32998 | 16032 | 205.83% |

| | | | |
|-------------|--------|----------|---------|
| 1024 | 153912 | 67257 | 228.84% |
| 2048 | 803120 | 249253 | 322.21% |
| | | Average: | 235.40% |

Strength Reduction

Finally, we can use strength reduction to further improve performance. We notice that *iidim* and *iidimjj* increase by *dim* and *1*, respectively, in each iteration of their respective loops. Rather than fully recalculating them each time, we can calculate their initial value outside of their loops and then increment them in their loops.

After:

```
int ii, jj, iDim, iidim, iidimjj, iiEnd, jjEnd, dimMinusOne, iPlusOne, jPlusOne, iMinusOne, jMinusOne, iiStart,
jjStart;
pixel_sum sum;
pixel current_pixel;
dimMinusOne = dim - 1;

for (i = 0; i < dim; i++) {
    iPlusOne = i + 1;
    iMinusOne = i - 1;
    iDim = i * dim;
    for (j = 0; j < dim; j++) {
        jPlusOne = j + 1;
        jMinusOne = j - 1;
        sum.red = sum.green = sum.blue = 0;
        sum.num = 0;
        iiStart = (iMinusOne > 0 ? iMinusOne : 0);
        iidim = iiStart * dim;
        iiEnd = (iPlusOne < dimMinusOne ? iPlusOne : dimMinusOne);
        for (ii = iiStart; ii <= iiEnd; ii++) {
            jjStart = (jMinusOne > 0 ? jMinusOne : 0);
            iidimjj = iidim + jjStart;
            jjEnd = (jPlusOne < dimMinusOne ? jPlusOne : dimMinusOne);
            for (jj = jjStart; jj <= jjEnd; jj++) {
                sum.red += (int)src[iidimjj].red;
                sum.green += (int)src[iidimjj].green;
                sum.blue += (int)src[iidimjj].blue;
                sum.num++;
                iidimjj++;
            }
            iidim += dim;
        }

        current_pixel.red = (unsigned short)(sum.red / sum.num);
        current_pixel.green = (unsigned short)(sum.green / sum.num);
        current_pixel.blue = (unsigned short)(sum.blue / sum.num);
        dst[(iDim + j)] = current_pixel;
    }
}
```

This resulted in a slight increase in performance from our previous optimization, bringing the overall performance gain to ~140%. This improvement is because addition is less costly than multiplication and just iteratively adding to something already calculated is also less costly than recalculating the whole value.

| Dimension | naive_smooth | my_smooth | Speedup |
|-------------|--------------|-----------|---------|
| 256 | 7164 | 3830 | 187.05% |
| 512 | 30933 | 15172 | 203.88% |
| 1024 | 148973 | 59943 | 248.52% |
| 2048 | 761234 | 236379 | 322.04% |
| | | Average: | 240.37% |

Lastly, we notice that the above loops require logic and branching to check if the current pixel is a corner (which only averages 4 pixels), an edge (which averages 6), or an inside pixel that averages 9 pixels. This results in more computations being done for each pixel. To avoid this, we can separate calculating averages for the corners, edges, and inside from each other. And, to improve memory locality and increase cache hits, we can do the sides in a clockwise direction (starting at the top left corner, then the top edge, then the top right corner, then right side, and so on). For the corners, we also notice that we divide by 4 to calculate the average, so we can do a right shift of 2 instead since multiplying is more costly than a bit-shift. We can also use dataflow analysis to reduce common subexpressions and reduce the number of calculations. Putting all of these together, the code now is much longer but significantly more efficient.

After:

```
int dim2, dimMinusOne, dimPlusOne, iMinusOne, iPlusOne, k, l, m, n, o, p, q, r, s;
dim2 = dim * dim;
dimMinusOne = dim - 1;
dimPlusOne = dim + 1;
pixel current_pixel;

// top left corner
current_pixel.red = (src[0].red + src[1].red + src[dim].red + src[dimPlusOne].red)
>> 2;
current_pixel.green = (src[0].green + src[1].green + src[dim].green +
src[dimPlusOne].green) >> 2;
current_pixel.blue = (src[0].blue + src[1].blue + src[dim].blue +
src[dimPlusOne].blue) >> 2;
dst[0] = current_pixel;

// top side
iMinusOne = -1;
iPlusOne = 1;
l = dim;
k = dim - 1;
```

```

    j = dim + 1;
    for (i = 1; i < dimMinusOne; i++) {
        iMinusOne++;
        iPlusOne++;
        l++;
        k++;
        j++;
        current_pixel.red = (src[iMinusOne].red + src[i].red + src[iPlusOne].red +
src[k].red + src[l].red + src[j].red) / 6;
        current_pixel.green = (src[iMinusOne].green + src[i].green +
src[iPlusOne].green + src[k].green + src[l].green + src[j].green) / 6;
        current_pixel.blue = (src[iMinusOne].blue + src[i].blue + src[iPlusOne].blue +
src[k].blue + src[l].blue + src[j].blue) / 6;
        dst[i] = current_pixel;
    }

    // top right corner
    j = dimMinusOne + dim;
    k = j - 1;
    iMinusOne++;
    current_pixel.red = (src[iMinusOne].red + src[dimMinusOne].red + src[k].red +
src[j].red) >> 2;
    current_pixel.green = (src[iMinusOne].green + src[dimMinusOne].green + src[k].green
+ src[j].green) >> 2;
    current_pixel.blue = (src[iMinusOne].blue + src[dimMinusOne].blue + src[k].blue +
src[j].blue) >> 2;
    dst[dimMinusOne] = current_pixel;

    // right side
    k = dim2 - 1;
    n = j - 1;
    l = n - dim;
    m = j - dim;
    p = j + dim;
    o = p - 1;
    for (i = j; i < k; i += dim) {
        current_pixel.red = (src[l].red + src[m].red + src[n].red + src[i].red +
src[o].red + src[p].red) / 6;
        current_pixel.green = (src[l].green + src[m].green + src[n].green +
src[i].green + src[o].green + src[p].green) / 6;
        current_pixel.blue = (src[l].blue + src[m].blue + src[n].blue + src[i].blue +
src[o].blue + src[p].blue) / 6;
    }

```

```

        dst[i] = current_pixel;
        l += dim;
        n += dim;
        m += dim;
        o += dim;
        p += dim;
    }

    // bottom right corner
    l = k - 1;
    m = k - dim;
    n = m - 1;
    current_pixel.red = (src[n].red + src[m].red + src[l].red + src[k].red) >> 2;
    current_pixel.green = (src[n].green + src[m].green + src[l].green + src[k].green)
>> 2;
    current_pixel.blue = (src[n].blue + src[m].blue + src[l].blue + src[k].blue) >> 2;
    dst[k] = current_pixel;

    // bottom side, from right
    j = m + 1;
    m = l - dim;
    k = m - 1;
    n = m + 1;
    o = l - 1;
    p = l + 1;
    for (i = l; i > j; i--) {
        current_pixel.red = (src[k].red + src[m].red + src[n].red + src[o].red +
src[i].red + src[p].red) / 6;
        current_pixel.green = (src[k].green + src[m].green + src[n].green +
src[o].green + src[i].green + src[p].green) / 6;
        current_pixel.blue = (src[k].blue + src[m].blue + src[n].blue + src[o].blue +
src[i].blue + src[p].blue) / 6;
        dst[i] = current_pixel;
        m--;
        k--;
        n--;
        o--;
        p--;
    }

    // bottom left corner
    j = i - dim;

```

```

    k = j + 1;
    l = i + 1;
    current_pixel.red = (src[j].red + src[k].red + src[i].red + src[l].red) >> 2;
    current_pixel.green = (src[j].green + src[k].green + src[i].green + src[l].green)
>> 2;
    current_pixel.blue = (src[j].blue + src[k].blue + src[i].blue + src[l].blue) >> 2;
    dst[i] = current_pixel;

    // left side, from bottom
    k = j - dim;
    l = k + 1;
    m = j + 1;
    n = j + dim;
    o = n + 1;
    for (i = j; i > 0; i -= dim) {
        current_pixel.red = (src[k].red + src[l].red + src[i].red + src[m].red +
src[n].red + src[o].red) / 6;
        current_pixel.green = (src[k].green + src[l].green + src[i].green +
src[m].green + src[n].green + src[o].green) / 6;
        current_pixel.blue = (src[k].blue + src[l].blue + src[i].blue + src[m].blue +
src[n].blue + src[o].blue) / 6;
        dst[i] = current_pixel;
        k -= dim;
        l -= dim;
        m -= dim;
        n -= dim;
        o -= dim;
    }

    // inside
    k = dim;
    l = k - 1;
    m = k + 1;
    n = l - dim;
    o = k - dim;
    p = o + 1;
    q = l + dim;
    r = q + 1;
    s = r + 1;
    for (i = 1; i < dimMinusOne; i++) {
        for (j = 1; j < dimMinusOne; j++) {
            k++;

```

```

        l++;
        m++;
        n++;
        o++;
        p++;
        q++;
        r++;
        s++;

        current_pixel.red = (src[l].red + src[k].red + src[m].red + src[n].red +
src[o].red + src[p].red + src[q].red + src[r].red + src[s].red) / 9;

        current_pixel.green = (src[l].green + src[k].green + src[m].green +
src[n].green + src[o].green + src[p].green + src[q].green + src[r].green +
src[s].green) / 9;

        current_pixel.blue = (src[l].blue + src[k].blue + src[m].blue + src[n].blue
+ src[o].blue + src[p].blue + src[q].blue + src[r].blue + src[s].blue) / 9;

        dst[k] = current_pixel;
    }
    k += 2;
    l += 2;
    m += 2;
    n += 2;
    o += 2;
    p += 2;
    q += 2;
    r += 2;
    s += 2;
}

```

Smooth now sees a performance improvement of over 350%, with the performance increasing as the size of the image increases:

| Dimension | naive_smooth | my_smooth | Speedup |
|-------------|--------------|-----------|---------|
| 256 | 6799 | 1940 | 350.46% |
| 512 | 30430 | 7879 | 386.22% |
| 1024 | 155057 | 34212 | 453.22% |
| 2048 | 834334 | 134984 | 618.10% |
| | | Average: | 452.00% |